FPGA

Inference Engine Structure

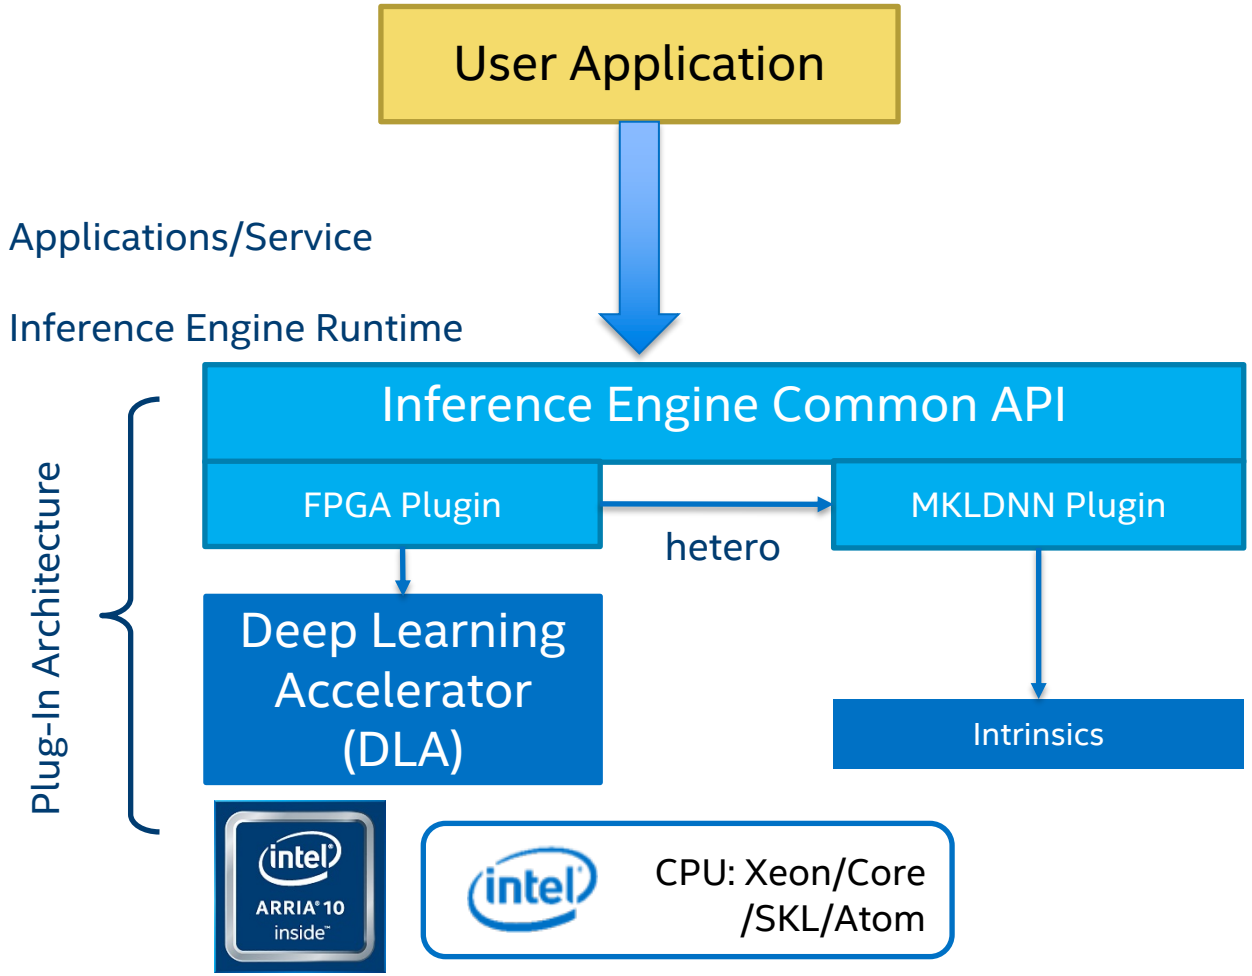

FPGA implementation is “Deep Learning Accelerator” (DLA)

Additional step: FPGA RTE (aocl) loads bitstream with desired version of DLA to FPGA before running

Supported layers

The following layers are supported by the plugin:

- Batch_norm (being converted by Model Optimizer to ScaleShift layer)
- Concat
- Convolution (dilated convolutions are supported, depthwise are not supported)
- Eltwise (operation sum is supported)
- Fully Connected
- LRN Normalization
- Pooling
- Power (scale and offset parameters are supported)
- ReLu (with negative slope)
- ScaleShift

Note Support is limited to the specific parameters (depending on the bitstream).

Heterogeneous execution

In case when topology contains layers not supported on FPGA, you need to use **Heterogeneous plugin** with dedicated fallback device.

Setup the FPGA Card

You have three options, more are coming up…

Intel® Arria® 10 GX FPGA Development Kit

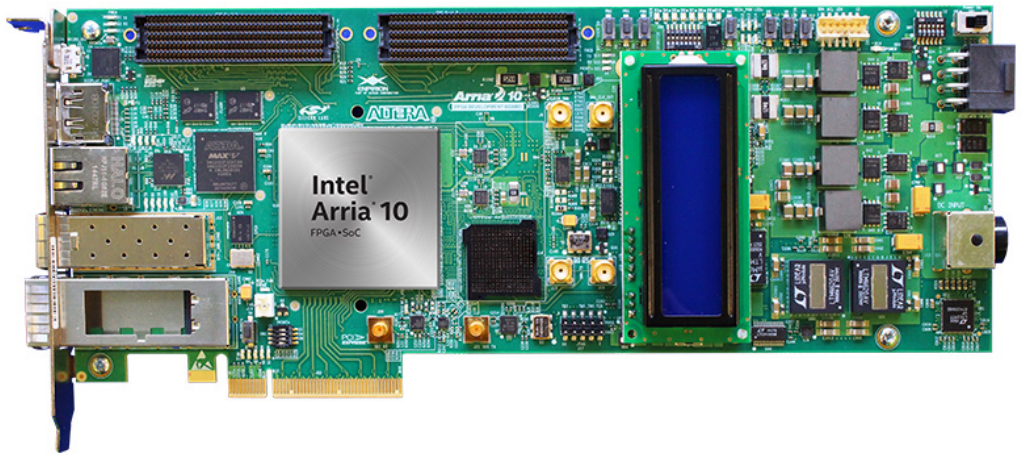

PCIe Gen3 x 8, full size card
(We are using this card in this session)

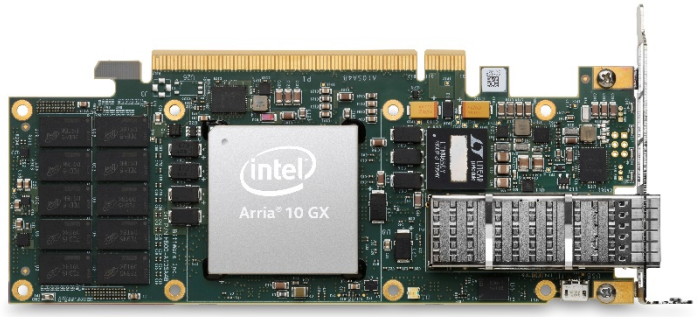

Intel® Programmable Acceleration Card with Intel Arria® 10 GX FPGA

PCIe Gen3 x 8, Half-height half-length card

Deep Learning Acceleration Card with Intel Arria® 10 GX FPGA

PCIe Gen3 x 8, Half-height half-length card

Setup the OS and Software

Setup a clean Ubuntu 16.04, upgrade the library and kernel to 4.13
Until it shows no further update is available…

```
$ sudo apt-get upgrade
[sudo] password for intel: xxxx
Reading package lists... Done
Building dependency tree
Reading state information... Done
Calculating upgrade... Done
0 upgraded, 0 newly installed, 0 to remove and 0 not upgraded.
$
```

Setup Quartus tools and OpenCL RT

Select, download and install Quartus Prime Pro Edition v17.1 Update 2

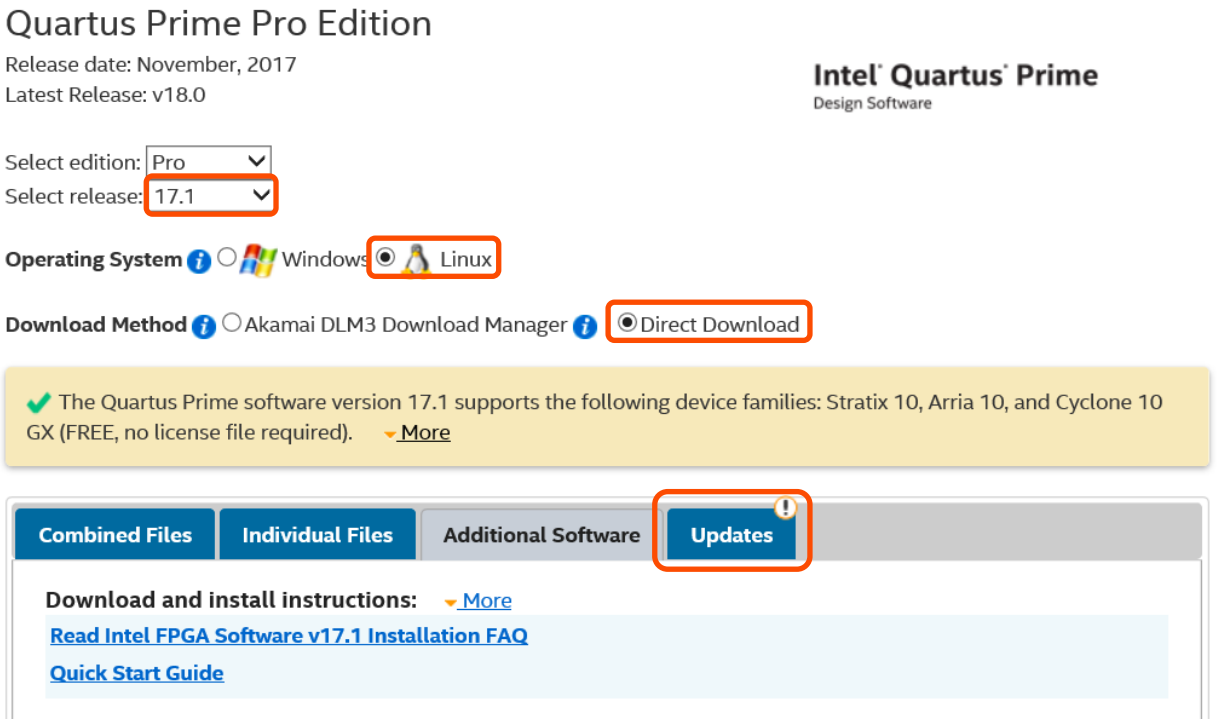

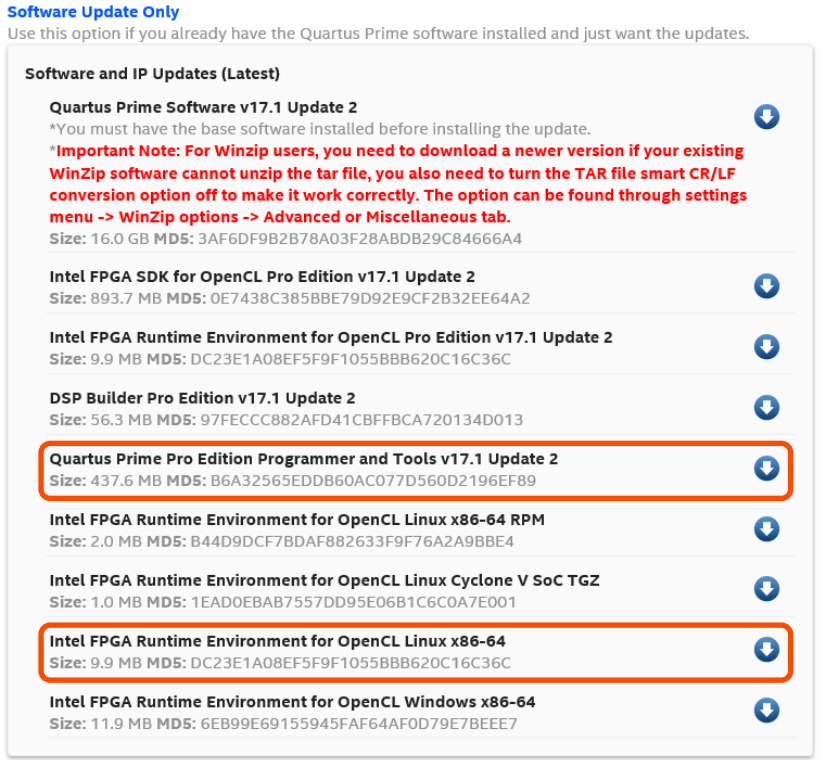

Setup the OS and Software

Initializing the FPGA card, and soft reboot…

```
$ export PATH=/opt/intelFPGA_pro/17.1/qprogrammer/bin:$PATH
$ quartus_pqm -c 1 -m JTAG -o "p;top.sof"
**********************************************************************
Running Quartus Prime Programmer
Version 17.1.2 Build 304 01/31/2018 SJ Pro Edition
...
Info: Quartus Prime Programmer was successful. 0 errors, 0 warnings
    Info: Peak virtual memory: 1628 megabytes
    Info: Processing ended: Fri May 11 08:07:51 2018
    Info: Elapsed time: 00:00:21
    Info: Total CPU time (on all processors): 00:00:10
$ sudo reboot
```

Setup the OS and Software

Setup OpenCL RT for FPGA…

```
$ export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:$INTELFPGAOCLSDKROOT/host/
  linux64/lib:$AOCL_BOARD_PACKAGE_ROOT/linux64/lib
$ aocl install
aocl install: Running install from
/tools/aclboardpkg/altera_a10pciedk/16.0/linux64/libexec
...
  CC      /tmp/opencl_driver_x6GjWS/aclpci_a10_ref_drv.mod.o
  LD [M]  /tmp/opencl_driver_x6GjWS/aclpci_a10_ref_drv.ko.unsigned
  NO SIGN [M] /tmp/opencl_driver_x6GjWS/aclpci_a10_ref_drv.ko
make: Leaving directory `/usr/src/kernels/2.6.32-358.el6.x86_64'
$
```

Diagnostics

After reboot, run diagnosis and it should show PASSED status

```
$ aocl diagnose
-----------------------------------------------------------
...
Vendor: Intel Corporation

Phys Dev Name Status  Information
acla10_ref0   Passed   Arria 10 Reference Platform (acla10_ref0)
                       PCIe dev_id = 2494, bus:slot.func = 04:00.00,
                       Gen3 x8
                       FPGA temperature = 53.3438 degrees C.

DIAGNOSTIC_PASSED
-----------------------------------------------------------
$
```

Install the OpenVINO™

Follow the instructions to install Open Visual Inference & Neural network Optimization (OpenVINO™)

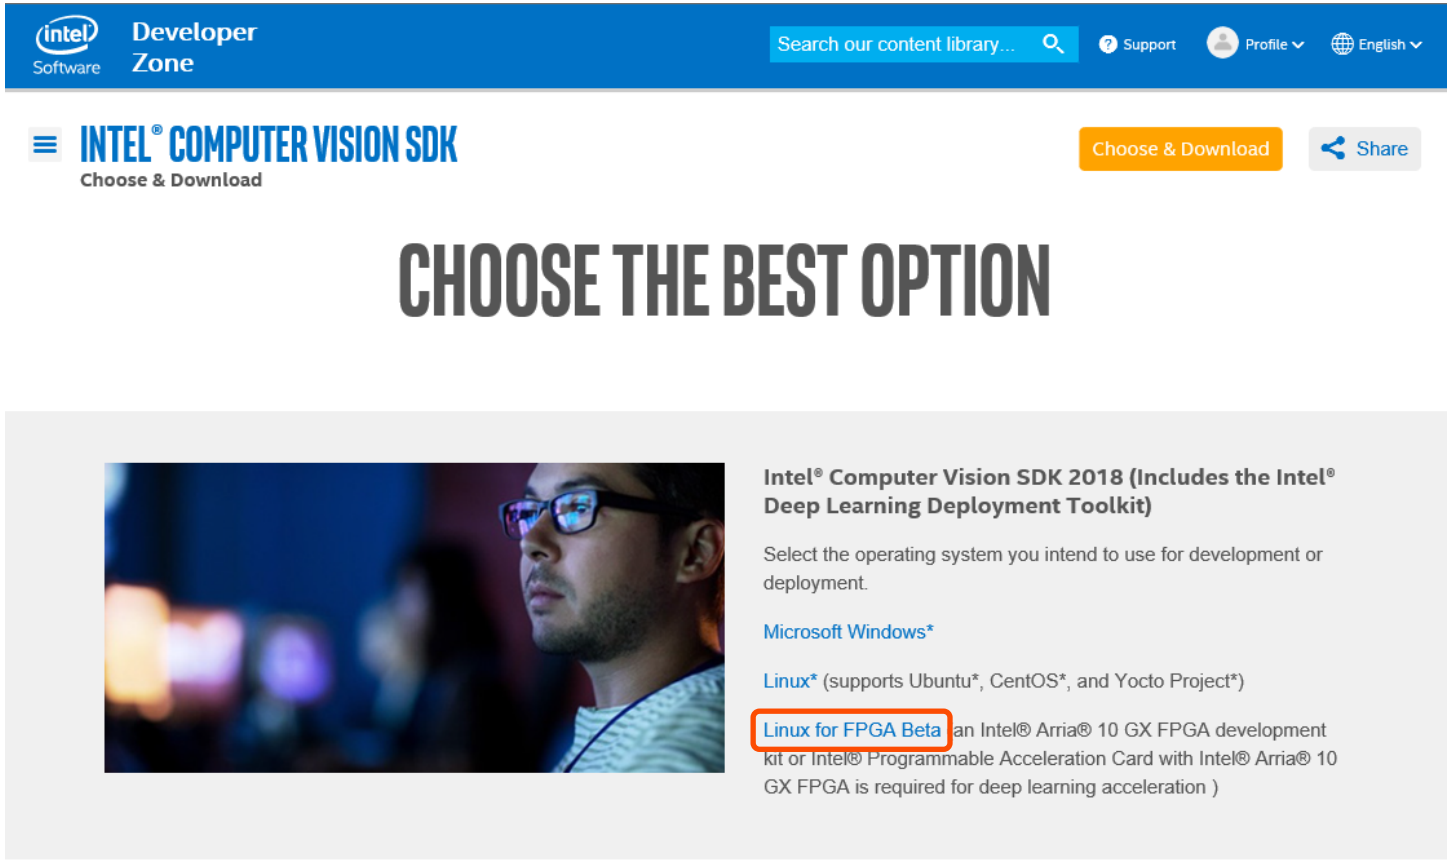

Run the Demo!

Targeting vehicle detection and classification workloads to CPU only

```
$ /opt/intel/computer_vision_sdk_fpga_2018.0.234/deployment_tools/inference_
  engine/samples/build/intel64/Release/security_barrier_camera_sample
  -d CPU
  -i /opt/intel/computer_vision_sdk_fpga_2018.0.234/deployment_tools/demo/
     test_video.mp4
  -m /opt/intel/computer_vision_sdk_fpga_2018.0.234/deployment_tools/demo/
     intel_models/vehicle-license-plate-detection-barrier-0007/FP32/vehicle-
     license-plate-detection-barrier-0007.xml
  -m_va opt/intel/computer_vision_sdk_fpga_2018.0.234/deployment_tools/demo/
        intel_models/vehicle-attributes-recognition-barrier-0010/FP32/
        vehicle-attributes-recognition-barrier-0010.xml
$
```

Setup the Environment

Before start deploying your workloads to FPGA, setup the environment…

```
$ source /opt/intel/computer_vision_sdk_fpga_2018.0.234/bin/setupvars.sh
$ export LD_LIBRARY_PATH=/opt/intel/computer_vision_sdk_fpga/deployment_
  tools/inference_engine/lib/ubuntu_16.04/intel64:$LD_LIBRARY_PATH
$ export LD_LIBRARY_PATH=/opt/altera/aocl-pro-rte/aclrte-linux64/host/
  linux64/lib:$LD_LIBRARY_PATH
$
```

Program the Bitstream to the FPGA

Setup bitstream variables, and load the bitstream to FPGA…

```
$ export CL_CONTEXT_COMPILER_MODE_INTELFPGA=3
$ export DLA_AOCX=/opt/intel/computer_vision_sdk_fpga_2018.0.234/a10_devkit_
  bitstreams/0-8-1_a10dk_fp16_8x48_arch06.aocx
$ source /opt/altera/aocl-pro-rte/aclrte-linux64/init_opencl.sh
$ aocl diagnose
$ aocl program acl0 $DLA_AOCX
Programming device: a10gx : Arria 10 Reference Platform (acla10_ref0)
Reprogramming device [0] with handle 1
Program succeed.
$
```

Inference Engine, few FPGA tips

- First iteration with FPGA is always significantly slower than the rest, make sure you run multiple iterations (“-ni”)
- If CPU utilization of concern, minimizes the busy wait time when OMP threads loop in between parallel regions:
  - set KMP_BLOCKTIME=0
  - consider playing OMP_NUM_THREADS
- FPGA performance heavily depends on the bitstream

Run the Demo in the Heterogeneous Mode

Setup the environment properly for heterogeneous mode

```
$ export KMP_BLOCKTIME=1
$ export OMP_NUM_THREADS=4
$
```

Run the Demo on FPGA!

Targeting vehicle detection and classification workloads to FPGA and CPU

```
$ /opt/intel/computer_vision_sdk_fpga_2018.0.234/deployment_tools/inference_
  engine/samples/build/intel64/Release/security_barrier_camera_sample
  -d HETERO:FPGA,CPU
  -i /opt/intel/computer_vision_sdk_fpga_2018.0.234/deployment_tools/demo/
     test_video.mp4
  -m /opt/intel/computer_vision_sdk_fpga_2018.0.234/deployment_tools/demo/
     intel_models/vehicle-license-plate-detection-barrier-0007/FP32/vehicle-
     license-plate-detection-barrier-0007.xml
  -m_va opt/intel/computer_vision_sdk_fpga_2018.0.234/deployment_tools/demo/
        intel_models/vehicle-attributes-recognition-barrier-0010/FP32/
        vehicle-attributes-recognition-barrier-0010.xml
$
```

Contrast

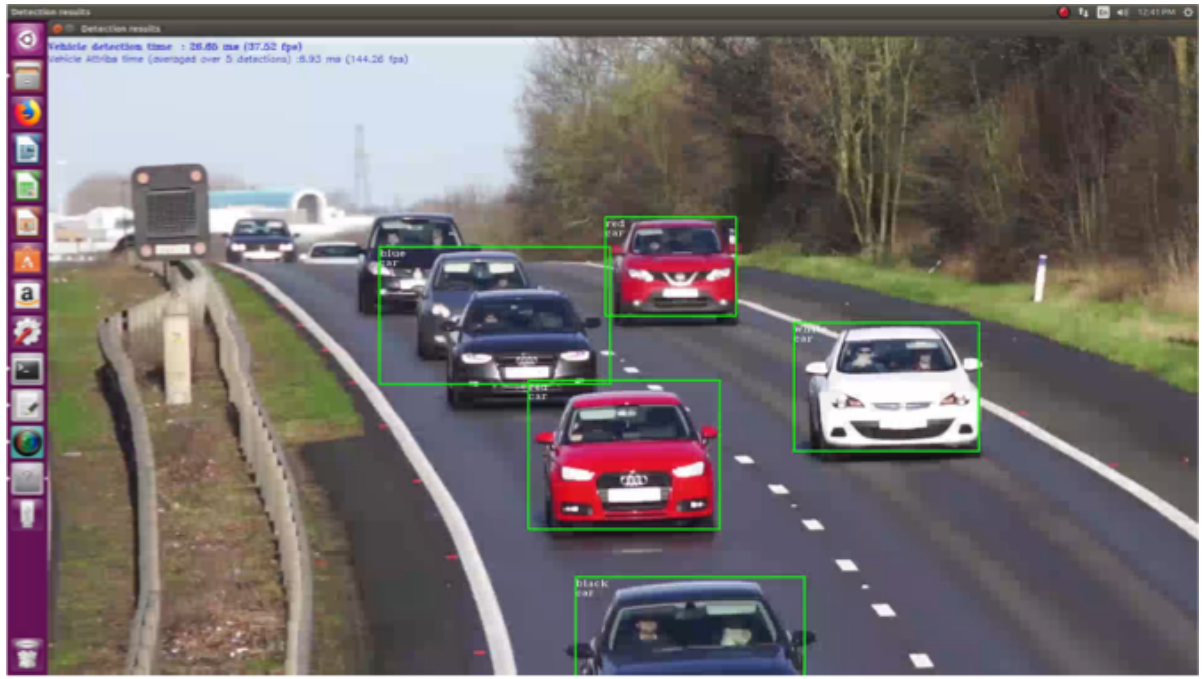

CPU only without OpenMP tweak

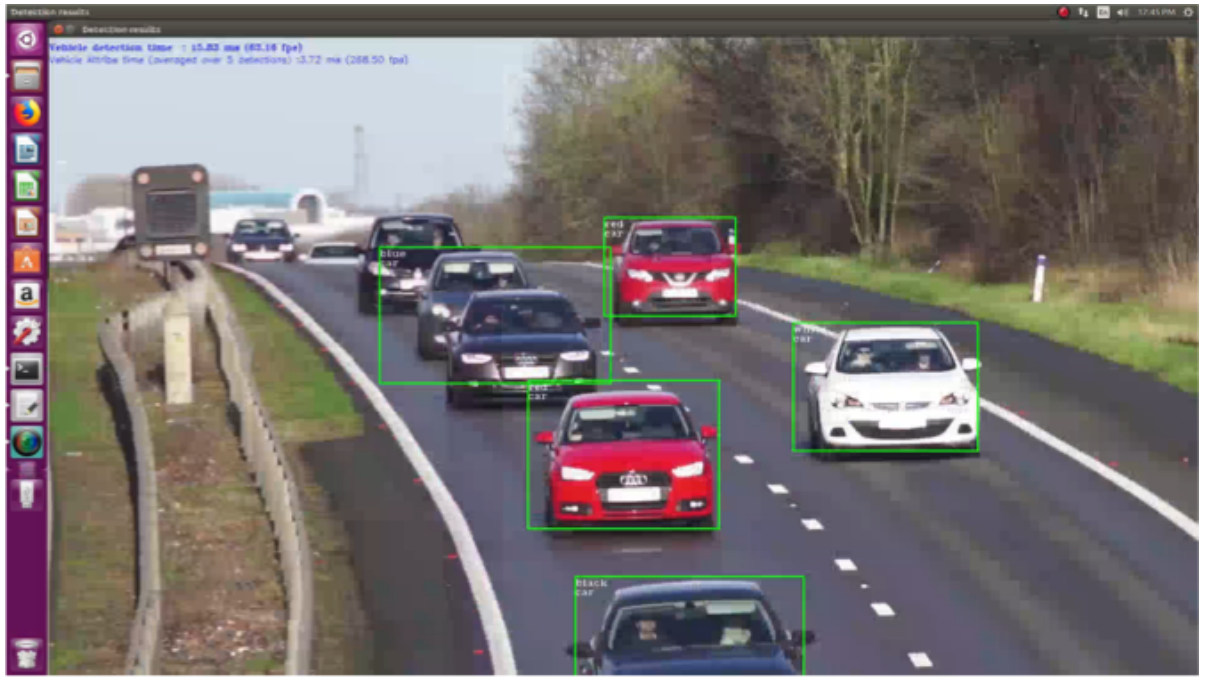

Heterogeneous mode (FPGA + CPU) with OpenMP tweak